

CPSC 351: Operating Systems

Project 2 — UNIX Shell

This project consists of designing a C program to serve as a shell interface that accepts user commands and then executes each command in a separate process. Your implementation will support input and output redirection, as well as pipes as a form of IPC between a pair of commands. Completing this project will involve using the UNIX **fork()**, **exec()**, **wait()**, **dup2()**, and **pipe()** system calls and can be completed on any Linux, UNIX, or macOS system.

I. Overview

A shell interface gives the user a prompt, after which the next command is entered. The example below illustrates the prompt **osh>** and the user's next command: **cat prog.c**. (This command displays the file **prog.c** on the terminal using the UNIX **cat** command.)

```
osh > cat prog.c
```

One technique for implementing a shell interface is to have the parent process first read what the user enters on the command line (in this case, **cat prog.c**) and then create a separate child process that performs the command. Unless otherwise specified, the parent process waits for the child to exit before continuing. This is similar in functionality to the new process creation illustrated in Figure 3.9. However, UNIX shells typically also allow the child process to run in the background, or concurrently. To accomplish this, we add an ampersand (&) at the end of the command. Thus, if we rewrite the above command as

```
osh > cat prog.c &
```

the parent and child processes will run concurrently.

The separate child process is created using the **fork()** system call, and the user's command is executed using one of the system calls in the **exec()** family (as described in Section 3.3.1).

A C program that provides the general operations of a command-line shell is supplied in Figure 3.36. The **main()** function presents the prompt **osh->** and outlines the steps to be taken after input from the user has been read. The **main()** function continually loops as long as **should run** equals 1; when the user enters **exit** at the prompt, your program will set **should run** to 0 and terminate.

```
#include <stdio.h>
#include <unistd.h>
#define MAX LINE 80 /* The maximum length command */

int main(void) {
    char *args[MAX LINE/2 + 1]; /* command line arguments */
    int should run = 1; /* flag to determine when to exit program */

    while (should run) {
        printf("osh>");
        fflush(stdout);

        /**
         * After reading user input, the steps are:
         * (1) fork a child process using fork()
         * (2) the child process will invoke execvp()
         * (3) parent will invoke wait() unless command included &
         */
    }
    return 0;
}
```

Figure 3.36 Outline of simple shell.

1. Creating the child process and executing the command in the child
2. Providing a history feature
3. Adding support of input and output redirection
4. Allowing the parent and child processes to communicate via a pipe
5. This project is organized into several parts:

II. Executing Command in a Child Process

The first task is to modify the **main()** function in Figure 3.36 so that a child process is forked and executes the command specified by the user. This will require parsing what the user has entered into separate tokens and storing the tokens in an array of character strings (args in Figure 3.36). For example, if the user enters the command

ps -ael at the osh > prompt,

the values stored in the args array are:

```
args[0] = "ps"
args[1] = "-ael"
args[2] = NULL
```

This args array will be passed to the **execvp()** function, which has the following prototype:

```
execvp(char *command, char *params[])
```

Here, command represents the command to be performed and params stores the parameters to this command. For this project, the **execvp()** function should be invoked as **execvp(args[0], args)**. Be sure to check whether the user included & to determine whether or not the parent process is to wait for the child to exit.

III. Creating a History Feature

The next task is to modify the shell interface program so that it provides a **history** feature to allow a user to execute the most recent command by entering **!!**. For example, if a user enters the command **ls -l**, she can then execute that command again by entering **!!** at the prompt. Any command executed in this fashion should be echoed on the user's screen, and the command should also be placed in the history buffer as the next command.

Your program should also manage basic error handling. If there is no recent command in the history, entering **!!** should result in a message **"No commands in history."**

IV. Redirecting Input and Output

Your shell should then be modified to support the '**>**' and '**<**' redirection operators, where '**>**' redirects the output of a command to a file and '**<**' redirects the input to a command from a file. For example, if a user enters

```
osh > ls > out.txt
```

the output from the ls command will be redirected to the file **out.txt**. Similarly, input can be redirected as well. For example, if the user enters

```
osh > sort < in.txt
```

the file in.txt will serve as input to the sort command.

Managing the redirection of both input and output will involve using the

dup2() function, which duplicates an existing file descriptor to another file descriptor. For example, if fd is a file descriptor to the file **out.txt**, the call

```
dup2(fd, STDOUT_FILENO);
```

duplicates **fd** to standard output (the terminal). This means that any writes to standard output will in fact be sent to the **out.txt** file.

You can assume that commands will contain either one input or one output redirection and will not contain both. In other words, you do not have to be concerned with command sequences such as **sort < in.txt > out.txt**.

V. Communication via a Pipe

The final modification to your shell is to allow the output of one command to serve as input to another using a pipe. For

example, the following command sequence

Project 2 -- Unix Shell page 2 of 3

osh > ls -l | less

has the output of the command **ls -l** serve as the input to the less command. Both the ls and less commands will run as separate processes and will communicate using the UNIX **pipe()** function described in Section 3.7.4. Perhaps the easiest way to create these separate processes is to have the parent process create the child process (which will execute **ls -l**). This child will also create another child process (which will execute less) and will establish a pipe between itself and the child process it creates. Implementing pipe functionality will also require using the **dup2()** function as described in the previous section. Finally, although several commands can be chained together using multiple pipes, you can assume that commands will contain only one pipe character and will not be combined with any redirection operators.

cpssc 351-Project Two --- Unix Shell --- due 25 Mar 2021 at 2359
Your name: Carson Carpenter
Repository: https://github.com/carsoncarpenter7/CS-351-Unix-Shell.git
Verify each of the following items and check the correct column. Incorrectly marked items will incur a 5% penalty.

Complete d	Not Complete d	Unix Shell
✓ <input type="checkbox"/>	<input type="checkbox"/>	Complete a Unix shell (osh) project that runs on any Linux, Unix or macOS system.
✓ <input type="checkbox"/>	<input type="checkbox"/>	Program is written in the C language, using fork() and exec() and wait()
✓ <input type="checkbox"/>	<input type="checkbox"/>	Shell program uses dup2() and pipe()
✓ <input type="checkbox"/>	<input type="checkbox"/>	Program creates a Unix shell with the parent process forking a child process
✓ <input type="checkbox"/>	<input type="checkbox"/>	Shell can run concurrently, e.g., osh> cat program.c &
✓ <input type="checkbox"/>	<input type="checkbox"/>	Shell can run in the background, e.g., osh> ls
✓ <input type="checkbox"/>	<input type="checkbox"/>	Shell terminates when exit() is typed, e.g., osh> exit()
✓ <input type="checkbox"/>	<input type="checkbox"/>	Shell allows multiple arguments to be passed through the command line
✓ <input type="checkbox"/>	<input type="checkbox"/>	Has a history feature that remembers the shell commands typed, e.g., osh>ls -l osh> !! will execute ls -l again
✓ <input type="checkbox"/>	<input type="checkbox"/>	Allows redirection of input using the dup2() function, and the < and > symbols osh> < input_file, or output, osh> > output_file, or both osh> < input_file > output_file

✓ <input type="checkbox"/>	<input type="checkbox"/>	Allows the parent and child process to communicate with a pipe (sends the output of ls to the less program) osh>ls -l less
Comments	Sometimes the shell does not run correctly when running concurrently	